


Homework 4

April 11, 2020

1 Recitation Exercises

1.1 Exercise 4

- a) Plot the probability of obtaining one point from each cluster in a sample of size K for values of K between 2 to 100.



- b) For K clusters, $K = 10, 100, \& 1000$, find the probability that a sample of size $2K$ contains at least one point from each cluster.

$$p = \frac{2K!}{K^K}$$

For $K = 10$,

$$p = \frac{2K!}{K^K} = \frac{2 \cdot 10!}{10^{10}} = 0.000728$$

For $K = 100$,

$$p = \frac{2K!}{K^K} = \frac{2 \cdot 100!}{100^{100}} = 1.867 \times 10^{-42}$$

For $K = 1000$,

$$p = \frac{2K!}{K^K} = \frac{2 \cdot 1000!}{1000^{1000}} = 0$$

1.2 Exercise 7

Given the data set: - there are m points & K clusters - half the points & clusters are in “more dense” regions - half the points & clusters are in “less dense” regions - the 2 regions are well-separated from each other

Which of the following should occur in order to minimize the squared error when finding K clusters:

The correct answer would be (c) which was to move centroids to the dense region. The less dense region would require more centroids if the squared error needs to be minimized. Recall that the less dense region tends to produce “noise” which would make it harder to identify clusters, hence needing more.

1.3 Exercise 11

Total SSE is the sum of the SSE for each separate attribute. - What does it mean if the SSE for one variable is low for all clusters?

If the SSE of one attribute is always low for all clusters, then the variable is just a constant.

- Low for just one cluster?

Then it be the opposite of the above, it would actually contribute to defining a cluster.

- High for all clusters?

If it's is high for every cluster, then I'd assume it is either noise or an outlier.

- High for one cluster?

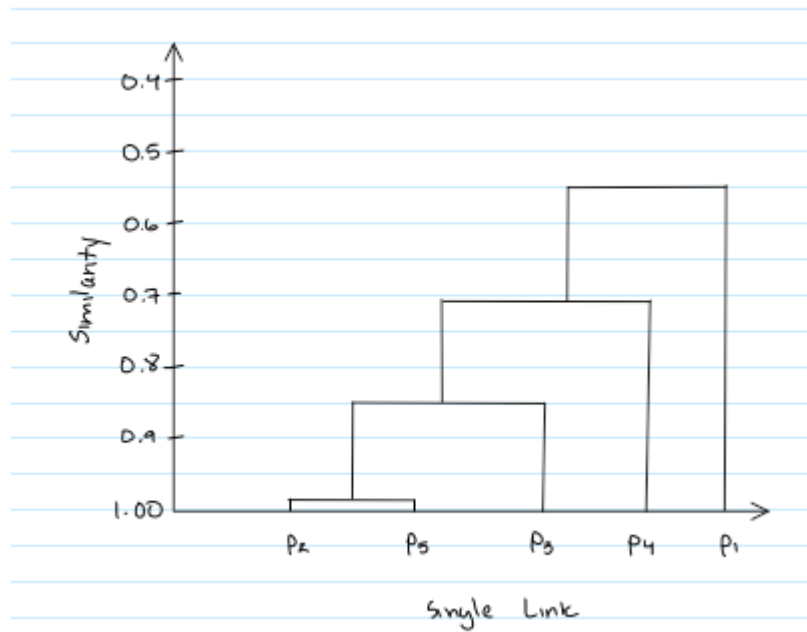
Another outlier, which would not help defining a cluster.

- How could you use the per variable SSE info to improve your clustering?


It can help in deciding which attributes to eliminate. Ch. 7 mentioned how sampling the data before clustering could be useful to eliminate the noise or outliers within the data, which would be useful to conserve time of the computation.

1.4 Exercise 16

- a) Single Link



b) Complete Link



1.5 Exercise 17

Given set of 1-dimensional points: $\{6, 12, 18, 24, 30, 42, 48\}$

- a) For each of the following sets of initial centroids, create 2 clusters by assigning each point to the nearest centroid & then calculate the total squared error for each of the 2 clusters.

i. $\{18, 45\}$

Reasoning: Using the minimal difference between points to find out which cluster they belong in... i.e $30 - 18 = 12$ vs $45 - 30 = 15$

Cluster 1 : $\{6, 12, 18, 24, 30\}$, Error = 360

Cluster 2: {42, 48}, Error = 18

Thus, total error = 378

ii. {15, 40}

Cluster 1 : {6, 12, 18, 24}, Error = 180

Cluster 2 : {30, 42, 48}, Error = 168

Thus, total error = 348

b) Do both sets of centroids represent stable solutions?

Yes, they do represent stable solutions since the above centroids represent centroids that are very far apart.

c) What are the 2 clusters produced by single link?

The minimal difference in (i) is $42-30 = 12$. The minimal difference in (ii) is $42-30 = 6$. So the two clusters formed by a single link is {6, 12, 18, 24, 30} & {42, 48}.

d) Which technique, K-means or single link, seems to produce the “most natural” clustering in this situation?

Since MIN usually produces the most natural clustering, I would go with MIN (single link).

e) What definition(s) of clustering does this natural clustering correspond to?

MIN produces contiguous clusters.

f) What well-known characteristic of the K-means alg. explains the previous behavior?

From what I recall, the K-means alg. is weak towards finding clusters that have a variety in sizes, or when not well-separated. The objective of minimizing squared error leads it to breaking the larger cluster, thus, producing the unnatural one in this case.

1.6 Exercise 21

Compute the entropy and purity for the confusion matrix...

Cluster #1:

$$\text{Entropy} = -\left[\left(\frac{1}{693}\right)\log\left(\frac{1}{693}\right) + \left(\frac{1}{693}\right)\log\left(\frac{1}{693}\right) + \left(\frac{0}{693}\right)\log\left(\frac{0}{693}\right) + \left(\frac{11}{693}\right)\log\left(\frac{11}{693}\right) + \left(\frac{4}{693}\right)\log\left(\frac{4}{693}\right) + \left(\frac{676}{693}\right)\log\left(\frac{676}{693}\right)\right] = 0.199 \approx 0.2, \text{Purity} = \frac{676}{693} = 0.975 \approx 0.98$$

Cluster #2 :

$$\text{Entropy} = 1.84, \text{Purity} = 0.53$$

Cluster #3:

$$\text{Entropy} = 1.7, \text{Purity} = 0.49$$

Total:

$$\text{Entropy} = 1.44, 0.61$$

1.7 Exercise 22

Given 2 sets of 100 points that fall within the unit square. One set of points is arranged so that the points are uniformly spaced. The other set of points is generated from a uniform distribution over the unit square.

- a) Is there a difference between the 2 set of points?

Definitely, the random points will have a region of less & more density, while the uniformly spaced will have uniform density.

- b) If so, which set of points will typically have a smaller SSE for $K=10$ clusters?

The random generated will have smaller SSE for $K=10$ clusters.

- c) What will be the behavior of DBSCAN on the uniform data set? The random data set?

Depending on the threshold, DBSCAN will either merge all the points in the uniform data set into a cluster or state they are all just noise. In terms of the random data set, DBSCAN can often find clusters in random data due to the variety of density between regions.

2 Practicum Problems

2.1 Problem 1 - Auto-Mpg Dataset

```
[163]: #All necessary imports
import pandas as pd
import numpy as np
from matplotlib import pyplot as plt
from sklearn.impute import SimpleImputer
#from sklearn.datasets.samples_generator import make_blobs
#from sklearn.cluster import KMeans
from sklearn.cluster import AgglomerativeClustering
import scipy.cluster.hierarchy as sch
```

```
[164]: #Load the auto-mpg sample dataset
autompg_ds = pd.read_csv('auto-mpg.csv', na_values=["?"])
autompg_ds = autompg_ds.drop(columns=['name', 'cylinders', 'model_year', '
→ 'origin'])
```

```
[165]: #Check to make sure the dataset works
autompg_ds.describe()
```

```
[165]:
```

	mpg	displacement	horsepower	weight	acceleration
count	398.000000	398.000000	392.000000	398.000000	398.000000
mean	23.514573	193.425879	104.469388	2970.424623	15.568090
std	7.815984	104.269838	38.491160	846.841774	2.757689
min	9.000000	68.000000	46.000000	1613.000000	8.000000
25%	17.500000	104.250000	75.000000	2223.750000	13.825000
50%	23.000000	148.500000	93.500000	2803.500000	15.500000

75%	29.000000	262.000000	126.000000	3608.000000	17.175000
max	46.600000	455.000000	230.000000	5140.000000	24.800000

```
[166]: #Impute any missing values with the mean of the dataset
imp_mean = SimpleImputer(missing_values=np.nan, strategy='mean')
imp_mean.fit(autompg_ds)
autompg_ds[autompg_ds.columns] = imp_mean.fit_transform(autompg_ds)
```

```
[167]: #Check the dataset again
autompg_ds.describe()
```

```
[167]:
```

	mpg	displacement	horsepower	weight	acceleration
count	398.000000	398.000000	398.000000	398.000000	398.000000
mean	23.514573	193.425879	104.469388	2970.424623	15.568090
std	7.815984	104.269838	38.199187	846.841774	2.757689
min	9.000000	68.000000	46.000000	1613.000000	8.000000
25%	17.500000	104.250000	76.000000	2223.750000	13.825000
50%	23.000000	148.500000	95.000000	2803.500000	15.500000
75%	29.000000	262.000000	125.000000	3608.000000	17.175000
max	46.600000	455.000000	230.000000	5140.000000	24.800000

```
[168]: #Perform Hierarchial Clustering with linkage set to average & default affinity
        ↳set to a euclidean.
#Remaining parameters must obtain a shallow tree with 3 clusters as targets

clustering = AgglomerativeClustering(n_clusters=3, affinity='euclidean',
        ↳linkage='average').fit(autompg_ds)
```

```
[169]: labels = clustering.labels_
print(labels)
```

```
[2 2 2 2 2 1 1 1 1 2 2 2 2 0 0 0 0 0 0 0 0 0 0 0 0 1 1 1 1 0 0 0 0 0 2 2 0
0 1 1 1 1 1 1 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 1 1 2 1 1 1 1 0 2 1
1 1 0 0 0 0 0 0 0 0 0 0 1 2 2 1 2 1 1 1 1 1 2 0 0 0 0 0 0 1 1 1 1 0 0 0 0
0 0 0 0 1 1 0 0 0 0 2 0 0 2 0 0 0 2 0 0 0 0 2 2 2 1 1 1 1 1 0 0 0 0 0 0 0
0 0 0 0 0 2 2 0 1 1 1 1 2 2 2 2 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 1 1 2 1 0 2 0 0 0 0 0 0 2 2 2 0 0 0 0 0 0 2 0 0 2 1 1 2 2 0 0 0 0 0 2
1 1 1 2 2 2 2 1 1 1 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 2 2 2 2 0 0 0 2 0 2
0 2 2 2 2 0 1 0 0 0 0 0 0 0 0 0 0 0 2 0 0 0 0 0 0 2 2 2 2 2 1 1 2 2 0 0 0
0 2 2 0 2 0 0 0 0 0 0 0 0 0 0 0 0 0 2 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 2 2 0 2 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
```

```
[170]: #To find mean and variance of each cluster, let's place the labels into the
        ↳dataframe which range from 0 to 2
autompg_ds['cluster'] = clustering.labels_
autompg_ds.describe()
```

```
[170]:
```

	mpg	displacement	horsepower	weight	acceleration \
count	398.000000	398.000000	398.000000	398.000000	398.000000
mean	23.514573	193.425879	104.469388	2970.424623	15.568090
std	7.815984	104.269838	38.199187	846.841774	2.757689
min	9.000000	68.000000	46.000000	1613.000000	8.000000
25%	17.500000	104.250000	76.000000	2223.750000	13.825000
50%	23.000000	148.500000	95.000000	2803.500000	15.500000
75%	29.000000	262.000000	125.000000	3608.000000	17.175000
max	46.600000	455.000000	230.000000	5140.000000	24.800000

	cluster
count	398.000000
mean	0.502513
std	0.770190
min	0.000000
25%	0.000000
50%	0.000000
75%	1.000000
max	2.000000

```
[171]: #This is the description for the first cluster
c0 = automp_g_ds.loc[automp_g_ds['cluster'] == 0]
c0.describe()
```

```
[171]:
```

	mpg	displacement	horsepower	weight	acceleration \
count	266.000000	266.000000	266.000000	266.000000	266.000000
mean	27.365414	131.934211	84.300061	2459.511278	16.298120
std	6.478913	53.179727	19.213107	427.354771	2.391296
min	13.000000	68.000000	46.000000	1613.000000	10.000000
25%	22.075000	97.000000	70.000000	2124.250000	14.550000
50%	26.900000	119.000000	85.000000	2395.000000	16.000000
75%	32.000000	149.750000	95.000000	2805.250000	17.600000
max	46.600000	455.000000	225.000000	3302.000000	24.800000

	cluster
count	266.0
mean	0.0
std	0.0
min	0.0
25%	0.0
50%	0.0
75%	0.0
max	0.0

```
[172]: #This is the description for the second cluster
c1 = automp_g_ds.loc[automp_g_ds['cluster'] == 1]
c1.describe()
```

```
[172]:
```

	mpg	displacement	horsepower	weight	acceleration	cluster
count	64.000000	64.000000	64.000000	64.000000	64.000000	64.0
mean	13.889062	358.093750	167.046875	4398.593750	13.025000	1.0
std	1.832781	46.240818	27.504937	272.602899	1.895106	0.0
min	9.000000	260.000000	110.000000	4042.000000	8.500000	1.0
25%	13.000000	318.000000	149.750000	4183.750000	12.000000	1.0
50%	14.000000	350.000000	156.500000	4357.000000	13.000000	1.0
75%	15.125000	400.000000	180.000000	4530.250000	14.000000	1.0
max	17.500000	455.000000	230.000000	5140.000000	19.000000	1.0

```
[173]: #This is the description for the third cluster
c2 = autompg_ds.loc[autompg_ds['cluster'] == 2]
c2.describe()
```

```
[173]:
```

	mpg	displacement	horsepower	weight	acceleration	cluster
count	68.000000	68.000000	68.000000	68.000000	68.000000	68.0
mean	17.510294	278.985294	124.470588	3624.838235	15.105882	2.0
std	2.971513	53.688847	26.703720	194.359999	3.249151	0.0
min	11.000000	163.000000	72.000000	3329.000000	8.000000	2.0
25%	15.000000	231.000000	105.000000	3435.250000	12.725000	2.0
50%	17.550000	260.000000	120.000000	3616.500000	15.450000	2.0
75%	19.125000	318.000000	150.000000	3782.000000	17.250000	2.0
max	26.600000	400.000000	190.000000	3988.000000	22.200000	2.0

```
[174]: #To find a difference if we had use the feature 'origin' as a class label,
↳let's create another dataset
autompg_ds2 = pd.read_csv('auto-mpg.csv', na_values=["?"])
autompg_ds2 = autompg_ds2.drop(columns=['name', 'cylinders', 'model_year'])
autompg_ds2.describe()
```

```
[174]:
```

	mpg	displacement	horsepower	weight	acceleration	\
count	398.000000	398.000000	392.000000	398.000000	398.000000	
mean	23.514573	193.425879	104.469388	2970.424623	15.568090	
std	7.815984	104.269838	38.491160	846.841774	2.757689	
min	9.000000	68.000000	46.000000	1613.000000	8.000000	
25%	17.500000	104.250000	75.000000	2223.750000	13.825000	
50%	23.000000	148.500000	93.500000	2803.500000	15.500000	
75%	29.000000	262.000000	126.000000	3608.000000	17.175000	
max	46.600000	455.000000	230.000000	5140.000000	24.800000	

	origin
count	398.000000
mean	1.572864
std	0.802055
min	1.000000
25%	1.000000
50%	1.000000


```
75%      2.000000
max       3.000000
```

```
[175]: #Impute any missing values with the mean of the dataset
imp_mean = SimpleImputer(missing_values=np.nan, strategy='mean')
imp_mean.fit(autompg_ds2)
autompg_ds2[autompg_ds2.columns] = imp_mean.fit_transform(autompg_ds2)
```

```
[176]: #This is the description for the second cluster
o1 = autompg_ds2.loc[autompg_ds2['origin'] == 1]
o1.describe()
```

```
[176]:
```

	mpg	displacement	horsepower	weight	acceleration	origin
count	249.000000	249.000000	249.000000	249.000000	249.000000	249.0
mean	20.083534	245.901606	118.814769	3361.931727	15.033735	1.0
std	6.402892	98.501839	39.617323	794.792506	2.751112	0.0
min	9.000000	85.000000	52.000000	1800.000000	8.000000	1.0
25%	15.000000	151.000000	88.000000	2720.000000	13.000000	1.0
50%	18.500000	250.000000	105.000000	3365.000000	15.000000	1.0
75%	24.000000	318.000000	150.000000	4054.000000	16.900000	1.0
max	39.000000	455.000000	230.000000	5140.000000	22.200000	1.0

```
[177]: #This is the description for the third cluster
o2 = autompg_ds2.loc[autompg_ds2['origin'] == 2]
o2.describe()
```

```
[177]:
```

	mpg	displacement	horsepower	weight	acceleration	origin
count	70.000000	70.000000	70.000000	70.000000	70.000000	70.0
mean	27.891429	109.142857	81.241983	2423.300000	16.787143	2.0
std	6.723930	22.582079	20.264743	490.043191	3.045687	0.0
min	16.200000	68.000000	46.000000	1825.000000	12.200000	2.0
25%	24.000000	92.250000	70.000000	2067.250000	14.500000	2.0
50%	26.500000	104.500000	77.500000	2240.000000	15.700000	2.0
75%	30.650000	121.000000	90.750000	2769.750000	18.900000	2.0
max	44.300000	183.000000	133.000000	3820.000000	24.800000	2.0

```
[178]: #This is the description for the third cluster
o3 = autompg_ds2.loc[autompg_ds2['origin'] == 3]
o3.describe()
```

```
[178]:
```

	mpg	displacement	horsepower	weight	acceleration	origin
count	79.000000	79.000000	79.000000	79.000000	79.000000	79.0
mean	30.450633	102.708861	79.835443	2221.227848	16.172152	3.0
std	6.090048	23.140126	17.819199	320.497248	1.954937	0.0
min	18.000000	70.000000	52.000000	1613.000000	11.400000	3.0
25%	25.700000	86.000000	67.000000	1985.000000	14.600000	3.0
50%	31.600000	97.000000	75.000000	2155.000000	16.400000	3.0

75%	34.050000	119.000000	95.000000	2412.500000	17.550000	3.0
max	46.600000	168.000000	132.000000	2930.000000	21.000000	3.0

To answer the final question, I do not see a clear relationship between the cluster assignment and the class labels...The data seems to be pretty far apart or barely different in different attributes when comparing.

2.2 Problem 2 - Boston Dataset

```
[101]: #Imports
from sklearn.cluster import KMeans
from sklearn.datasets import load_boston
from sklearn import preprocessing
from sklearn import metrics
```

```
[104]: #Load the dataset into a dataframe
boston_ds = load_boston()
boston_df = pd.DataFrame(boston_ds.data, columns=boston_ds.feature_names)
boston_df.describe()
```

```
[104]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM \
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000

	AGE	DIS	RAD	TAX	PTRATIO	B \
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	68.574901	3.795043	9.549407	408.237154	18.455534	356.674032
std	28.148861	2.105710	8.707259	168.537116	2.164946	91.294864
min	2.900000	1.129600	1.000000	187.000000	12.600000	0.320000
25%	45.025000	2.100175	4.000000	279.000000	17.400000	375.377500
50%	77.500000	3.207450	5.000000	330.000000	19.050000	391.440000
75%	94.075000	5.188425	24.000000	666.000000	20.200000	396.225000
max	100.000000	12.126500	24.000000	711.000000	22.000000	396.900000

	LSTAT
count	506.000000
mean	12.653063
std	7.141062
min	1.730000
25%	6.950000
50%	11.360000

```
75%      16.955000
max       37.970000
```

```
[126]: #Scale the dataset
df_scaled = pd.DataFrame(data=preprocessing.scale(boston_ds.data),
    ↪ columns=boston_ds.feature_names)
df_scaled.describe()
```

```
[126]:
```

	CRIM	ZN	INDUS	CHAS	NOX \
count	5.060000e+02	5.060000e+02	5.060000e+02	5.060000e+02	5.060000e+02
mean	-8.787437e-17	-6.343191e-16	-2.682911e-15	4.701992e-16	2.490322e-15
std	1.000990e+00	1.000990e+00	1.000990e+00	1.000990e+00	1.000990e+00
min	-4.197819e-01	-4.877224e-01	-1.557842e+00	-2.725986e-01	-1.465882e+00
25%	-4.109696e-01	-4.877224e-01	-8.676906e-01	-2.725986e-01	-9.130288e-01
50%	-3.906665e-01	-4.877224e-01	-2.110985e-01	-2.725986e-01	-1.442174e-01
75%	7.396560e-03	4.877224e-02	1.015999e+00	-2.725986e-01	5.986790e-01
max	9.933931e+00	3.804234e+00	2.422565e+00	3.668398e+00	2.732346e+00

	RM	AGE	DIS	RAD	TAX \
count	5.060000e+02	5.060000e+02	5.060000e+02	5.060000e+02	5.060000e+02
mean	-1.145230e-14	-1.407855e-15	9.210902e-16	5.441409e-16	-8.868619e-16
std	1.000990e+00	1.000990e+00	1.000990e+00	1.000990e+00	1.000990e+00
min	-3.880249e+00	-2.335437e+00	-1.267069e+00	-9.828429e-01	-1.313990e+00
25%	-5.686303e-01	-8.374480e-01	-8.056878e-01	-6.379618e-01	-7.675760e-01
50%	-1.084655e-01	3.173816e-01	-2.793234e-01	-5.230014e-01	-4.646726e-01
75%	4.827678e-01	9.067981e-01	6.623709e-01	1.661245e+00	1.530926e+00
max	3.555044e+00	1.117494e+00	3.960518e+00	1.661245e+00	1.798194e+00

	PTRATIO	B	LSTAT
count	5.060000e+02	5.060000e+02	5.060000e+02
mean	-9.205636e-15	8.163101e-15	-3.370163e-16
std	1.000990e+00	1.000990e+00	1.000990e+00
min	-2.707379e+00	-3.907193e+00	-1.531127e+00
25%	-4.880391e-01	2.050715e-01	-7.994200e-01
50%	2.748590e-01	3.811865e-01	-1.812536e-01
75%	8.065758e-01	4.336510e-01	6.030188e-01
max	1.638828e+00	4.410519e-01	3.548771e+00

```
[112]: #Perform K-Means with the scaled data & number of clusters = 2
clust_model = KMeans(n_clusters=2, init='k-means++')
clust_labels = clust_model.fit_predict(df_scaled)
print(clust_labels)
```

```
[1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
1 1 1 1 1 1 1 1 1 1 1 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
```

```
[115]: #Provide the Silhouette
silhouette_avg = metrics.silhouette_score(df_scaled, clust_labels)
print(silhouette_avg)
```

```
[116]: #Perform K-Means with the scaled data & number of clusters = 3
        clust_model = KMeans(n_clusters=3, init='k-means++')
        clust_labels = clust_model.fit_predict(df_scaled)
        print(clust_labels)
```

```
[117]: #Provide the Silhouette
silhouette_avg = metrics.silhouette_score(df_scaled, clust_labels)
print(silhouette_avg)
```

```
[118]: #Perform K-Means with the scaled data & number of clusters = 4
       clust_model = KMeans(n_clusters=4, init='k-means++')
       clust_labels = clust_model.fit_predict(df_scaled)
       print(clust_labels)
```

```
[119]: #Provide the Silhouette
silhouette_avg = metrics.silhouette_score(df_scaled, clust_labels)
print(silhouette_avg)
```

```
[122]: #Perform K-Means with the scaled data & number of clusters = 5  
clust_model = KMeans(n_clusters=5, init='k-means++')  
clust_labels = clust_model.fit_predict(df_scaled)  
print(clust_labels)
```

```
[123]: #Provide the Silhouette
silhouette_avg = metrics.silhouette_score(df_scaled, clust_labels)
print(silhouette_avg)
```

```
[124]: #Perform K-Means with the scaled data & number of clusters = 6
clust_model = KMeans(n_clusters=6, init='k-means++')
clust_labels = clust_model.fit_predict(df_scaled)
print(clust_labels)
```

```
[4 4 4 4 4 4 4 4 5 4 4 4 4 4 5 4 4 5 4 5 5 5 5 5 5 5 5 5 5 5 4 4
4 4 1 1 4 4 4 4 4 4 4 5 4 4 4 4 4 1 1 1 1 4 4 4 4 4 4 4 1 1 4 4 4 4 4 4
4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 5 4 5 5 5 5 5 5 5 5
5 5 5 5 5 5 5 5 5 5 5 5 5 5 5 5 5 5 5 5 5 5 5 5 5 5 5 5 2 5 5 5 5 5
5 5 5 5 2 5 2 2 5 5 5 5 2 5 2 2 5 5 5 5 5 5 5 5 4 4 4 4 4 4 4 4 4 4 4
4 4 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 4 4 5 2 2 2 2 2 4 5 4 2 5 2 2 2 2
2 4 4 4 4 4 4 4 4 4 4 4 2 4 2 4 1 4 4 4 4 1 4 4 4 4 4 4 4 4 1 1 1 1 1 4 4
4 4 4 4 4 4 4 4 4 4 2 4 4 4 2 2 4 2 2 4 4 4 4 2 1 1 1 1 1 1 1 1 1 4 4 4
4 4 1 1 1 4 1 1 4 4 4 4 4 5 5 4 5 5 5 5 5 5 4 4 4 4 4 4 4 4 4 4 4 4 1
4 4 4 4 4 4 4 4 1 4 1 1 4 4 1 1 1 1 1 1 1 1 2 2 2 0 0 0 0 2 2 0 0 3 0 2
2 0 2 0 0 0 0 0 0 0 3 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 3 3 0
0 0 3 3 3 3 3 3 3 3 3 3 0 0 0 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 0 0 0 0 0
0 3 0 0 0 0 3 0 0 0 3 3 3 3 0 0 0 0 0 0 0 3 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 5 5 5 5 5 5 5 5 5 5 5 5 5 5]
```

```
[125]: #Provide the Silhouette
silhouette_avg = metrics.silhouette_score(df_scaled, clust_labels)
print(silhouette_avg)
```

0.2617743548302116

Since the highest silhouette score was when the total # of clusters of 2, it is the optimal k because when a high value is shown, it indicates that the object is well matched to its own cluster & poorly matched to neighboring clusters.

```
[131]: #Redo the cluster with n_clusters = 2 to calculate the mean of all values
clust_model = KMeans(n_clusters=2, init='k-means++')
clust_labels = clust_model.fit_predict(df_scaled)
print(clust_labels)
```

[illegible]

```
1 1 1 1 1 1 1 1 1 1 1 1 0 0 0 0 0 0 0 0 0 0 0]
```

```
[132]: #Add the labels to the df_scaled
df_scaled['cluster'] = clust_labels
df_scaled.describe()
```

```
[132]:
```

	CRIM	ZN	INDUS	CHAS	NOX \
count	5.060000e+02	5.060000e+02	5.060000e+02	5.060000e+02	5.060000e+02
mean	-8.787437e-17	-6.343191e-16	-2.682911e-15	4.701992e-16	2.490322e-15
std	1.000990e+00	1.000990e+00	1.000990e+00	1.000990e+00	1.000990e+00
min	-4.197819e-01	-4.877224e-01	-1.557842e+00	-2.725986e-01	-1.465882e+00
25%	-4.109696e-01	-4.877224e-01	-8.676906e-01	-2.725986e-01	-9.130288e-01
50%	-3.906665e-01	-4.877224e-01	-2.110985e-01	-2.725986e-01	-1.442174e-01
75%	7.396560e-03	4.877224e-02	1.015999e+00	-2.725986e-01	5.986790e-01
max	9.933931e+00	3.804234e+00	2.422565e+00	3.668398e+00	2.732346e+00

	RM	AGE	DIS	RAD	TAX \
count	5.060000e+02	5.060000e+02	5.060000e+02	5.060000e+02	5.060000e+02
mean	-1.145230e-14	-1.407855e-15	9.210902e-16	5.441409e-16	-8.868619e-16
std	1.000990e+00	1.000990e+00	1.000990e+00	1.000990e+00	1.000990e+00
min	-3.880249e+00	-2.335437e+00	-1.267069e+00	-9.828429e-01	-1.313990e+00
25%	-5.686303e-01	-8.374480e-01	-8.056878e-01	-6.379618e-01	-7.675760e-01
50%	-1.084655e-01	3.173816e-01	-2.793234e-01	-5.230014e-01	-4.646726e-01
75%	4.827678e-01	9.067981e-01	6.623709e-01	1.661245e+00	1.530926e+00
max	3.555044e+00	1.117494e+00	3.960518e+00	1.661245e+00	1.798194e+00

	PTRATIO	B	LSTAT	cluster
count	5.060000e+02	5.060000e+02	5.060000e+02	506.000000
mean	-9.205636e-15	8.163101e-15	-3.370163e-16	0.349802
std	1.000990e+00	1.000990e+00	1.000990e+00	0.477379
min	-2.707379e+00	-3.907193e+00	-1.531127e+00	0.000000
25%	-4.880391e-01	2.050715e-01	-7.994200e-01	0.000000
50%	2.748590e-01	3.811865e-01	-1.812536e-01	0.000000
75%	8.065758e-01	4.336510e-01	6.030188e-01	1.000000
max	1.638828e+00	4.410519e-01	3.548771e+00	1.000000

```
[134]: #Show the mean values for all features in the first cluster
c0 = df_scaled.loc[df_scaled['cluster'] == 0]
c0.mean()
```

```
[134]: CRIM      -0.390124
      ZN        0.262392
      INDUS    -0.620368
      CHAS      0.002912
      NOX      -0.584675
      RM       0.243315
      AGE     -0.435108
```

```
DIS      0.457222
RAD      -0.583801
TAX      -0.631460
PTRATIO  -0.285808
B        0.326451
LSTAT    -0.446421
cluster  0.000000
dtype: float64
```

```
[135]: #Show the mean values for all features in the second cluster
c1 = df_scaled.loc[df_scaled['cluster'] == 1]
c1.mean()
```

```
[135]: CRIM      0.725146
ZN        -0.487722
INDUS     1.153113
CHAS      -0.005412
NOX       1.086769
RM        -0.452263
AGE       0.808760
DIS       -0.849865
RAD       1.085145
TAX       1.173731
PTRATIO   0.531248
B        -0.606793
LSTAT     0.829787
cluster   1.000000
dtype: float64
```

To answer the final question for this problem, it is important to remember that the mean of a cluster is the same as the centroid coordinate.

2.3 Problem 3 - Wine Dataset

```
[138]: #Imports
from sklearn.datasets import load_wine
```

```
[139]: #Load the dataset
wine_ds = load_wine()
wine_scaled = pd.DataFrame(data=preprocessing.scale(wine_ds.data),
    ↪ columns=wine_ds.feature_names)
wine_scaled.describe()
```

```
[139]:
```

	alcohol	malic_acid	ash	alcalinity_of_ash \
count	1.780000e+02	1.780000e+02	1.780000e+02	1.780000e+02
mean	7.841418e-15	2.444986e-16	-4.059175e-15	-7.110417e-17
std	1.002821e+00	1.002821e+00	1.002821e+00	1.002821e+00

	magnesium	total_phenols	flavanoids	nonflavanoid_phenols
count	1.780000e+02	1.780000e+02	1.780000e+02	1.780000e+02
mean	-2.494883e-17	-1.955365e-16	9.443133e-16	-4.178929e-16
std	1.002821e+00	1.002821e+00	1.002821e+00	1.002821e+00
min	-2.088255e+00	-2.107246e+00	-1.695971e+00	-1.868234e+00
25%	-8.244151e-01	-8.854682e-01	-8.275393e-01	-7.401412e-01
50%	-1.222817e-01	9.595986e-02	1.061497e-01	-1.760948e-01
75%	5.096384e-01	8.089974e-01	8.490851e-01	6.095413e-01
max	4.371372e+00	2.539515e+00	3.062832e+00	2.402403e+00

	od280/od315_of_diluted_wines	proline
count	1.780000e+02	1.780000e+02
mean	2.126888e-15	-6.985673e-17
std	1.002821e+00	1.002821e+00
min	-1.895054e+00	-1.493188e+00
25%	-9.522483e-01	-7.846378e-01
50%	2.377348e-01	-2.337204e-01
75%	7.885875e-01	7.582494e-01
max	1.960915e+00	2.971473e+00

```
[1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1  
1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 0 2 0 0 0 0 0 0 0 0 0 1  
0 0 0 0 0 0 0 0 0 2 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0  
0 0 0 0 0 0 0 2 0 0 1 0 0 0 0 0 0 0 2 2 2 2 2 2 2 2 2 2 2 2 2  
2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2]
```

```
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0  
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0  
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0  
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0  
1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1]
```

0.8900387051166616 0.4745258086227089

To answer the final question, homogeneity means all of the observations with the same class label are within the same cluster. Completeness means all members of the same class are within the same cluster. Both scores range from 0 to 1, with the higher $\#$ being the best outcome.

$$[\]:$$